

```
import pandas as pd
import numpy as np
import seaborn as sns
from datetime import datetime
import matplotlib.pyplot as plt
sns.set_style("whitegrid")
```

```
# importing the dataset
ABNB="https://raw.githubusercontent.com/ukantjadia/30-days-of-Mahcine-Learning/Main/DAY-01/ABNB.csv"
df = pd.read_csv(ABNB)
```

```
# Summary Stats

df.describe()
```

| | Open | High | Low | Close | Adj Close | Volume |
|-------|------------|------------|------------|------------|------------|--------------|
| count | 454.000000 | 454.000000 | 454.000000 | 454.000000 | 454.000000 | 4.540000e+02 |
| mean | 152.735925 | 156.465595 | 148.969855 | 152.729185 | 152.729185 | 6.690328e+06 |
| std | 28.340932 | 29.107813 | 27.490223 | 28.226659 | 28.226659 | 4.417387e+06 |
| min | 88.879997 | 91.459999 | 86.709999 | 89.080002 | 89.080002 | 1.995400e+06 |
| 25% | 136.821507 | 139.842495 | 132.920002 | 135.955002 | 135.955002 | 4.114375e+06 |
| 50% | 154.932495 | 159.099998 | 151.230004 | 155.654999 | 155.654999 | 5.450000e+06 |
| 75% | 172.556507 | 175.974998 | 168.595005 | 171.744996 | 171.744996 | 7.631375e+06 |
| max | 216.240005 | 219.940002 | 209.089996 | 216.839996 | 216.839996 | 3.975500e+07 |





```
import plotly.graph_objs as go
fig = go.Figure([go.Ohlc(x=df['Date'],
                        open=df['Open'],
                        high=df['High'],
                        low=df['Low'],
                        close=df['Close'])])
fig.update(layout_xaxis_rangeslider_visible=True)
fig.show()
```



```
plt.figure(figsize=(30,15));
dff = ['Open','Close','High','Low']
for no,col in enumerate(dff,1):
    plt.subplot(2,2,no)
    df[str(col)].hist(bins=50)
    plt.xlabel('Daily Return')
```

```
plt.title(f'{col}')
```

```
plt.tight_layout()
```



```
# Creating the dataframe of date formate and setting it to index
df['Date'] = pd.to_datetime(df['Date'])
df2 = df.set_index('Date')
df2.head()
```

| | Open | High | Low | Close | Adj Close | Volume |
|------------|------------|------------|------------|------------|------------|----------|
| Date | | | | | | |
| 2020-12-11 | 146.550003 | 151.500000 | 135.100006 | 139.250000 | 139.250000 | 26980800 |
| 2020-12-14 | 135.000000 | 135.300003 | 125.160004 | 130.000000 | 130.000000 | 16966100 |
| 2020-12-15 | 126.690002 | 127.599998 | 121.500000 | 124.800003 | 124.800003 | 10914400 |
| 2020-12-16 | 125.830002 | 142.000000 | 124.910004 | 137.990005 | 137.990005 | 20409600 |
| 2020-12-17 | 143.000000 | 152.449997 | 142.669998 | 147.050003 | 147.050003 | 15054700 |

```
def create_features_datetime(df):
    df['Year']=df.index.year
    df['Month']=df.index.month
    df['dow']=df.index.day_of_week

    return df
df_tr=create_features_datetime(df2)
df_tr.head()
```

| | Open | High | Low | Close | Adj Close | Volume | Year | Month | dow |
|------------|------------|------------|------------|------------|------------|----------|------|-------|-----|
| Date | | | | | | | | | |
| 2020-12-11 | 146.550003 | 151.500000 | 135.100006 | 139.250000 | 139.250000 | 26980800 | 2020 | 12 | 4 |
| 2020-12-14 | 135.000000 | 135.300003 | 125.160004 | 130.000000 | 130.000000 | 16966100 | 2020 | 12 | 0 |
| 2020-12-15 | 126.690002 | 127.599998 | 121.500000 | 124.800003 | 124.800003 | 10914400 | 2020 | 12 | 1 |
| 2020-12-16 | 125.830002 | 142.000000 | 124.910004 | 137.990005 | 137.990005 | 20409600 | 2020 | 12 | 2 |
| 2020-12-17 | 143.000000 | 152.449997 | 142.669998 | 147.050003 | 147.050003 | 15054700 | 2020 | 12 | 3 |

```
X = df_tr.drop('Close',axis=1)
y = df_tr['Close']
print(X)
```

| | Open | High | Low | Adj Close | Volume | Year | \ |
|------------|------------|------------|------------|------------|----------|------|---|
| Date | | | | | | | |
| 2020-12-11 | 146.550003 | 151.500000 | 135.100006 | 139.250000 | 26980800 | 2020 | |
| 2020-12-14 | 135.000000 | 135.300003 | 125.160004 | 130.000000 | 16966100 | 2020 | |
| 2020-12-15 | 126.690002 | 127.599998 | 121.500000 | 124.800003 | 10914400 | 2020 | |
| 2020-12-16 | 125.830002 | 142.000000 | 124.910004 | 137.990005 | 20409600 | 2020 | |
| 2020-12-17 | 143.000000 | 152.449997 | 142.669998 | 147.050003 | 15054700 | 2020 | |
| ... | ... | ... | ... | ... | ... | ... | |
| 2022-09-26 | 101.779999 | 105.360001 | 101.559998 | 103.230003 | 5278800 | 2022 | |
| 2022-09-27 | 106.040001 | 108.169998 | 104.139999 | 106.370003 | 5081000 | 2022 | |
| 2022-09-28 | 106.700996 | 111.129997 | 105.779999 | 110.690002 | 4518200 | 2022 | |
| 2022-09-29 | 109.190002 | 109.474998 | 105.169998 | 106.660004 | 4259500 | 2022 | |
| 2022-09-30 | 104.760002 | 109.080002 | 104.730003 | 105.040001 | 4747800 | 2022 | |
| Month | dow | | | | | | |
| Date | | | | | | | |
| 2020-12-11 | 12 | 4 | | | | | |
| 2020-12-14 | 12 | 0 | | | | | |
| 2020-12-15 | 12 | 1 | | | | | |
| 2020-12-16 | 12 | 2 | | | | | |
| 2020-12-17 | 12 | 3 | | | | | |
| ... | ... | ... | | | | | |
| 2022-09-26 | 9 | 0 | | | | | |
| 2022-09-27 | 9 | 1 | | | | | |
| 2022-09-28 | 9 | 2 | | | | | |
| 2022-09-29 | 9 | 3 | | | | | |
| 2022-09-30 | 9 | 4 | | | | | |

[454 rows x 8 columns]

```
# test size 20% and train size 80%

from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.2,random_state=0)
X_train, y_train = np.array(X_train), np.array(y_train)
X_test, y_test = np.array(X_test),np.array(y_test)
print(f"{y_test.shape} {X_test.shape}")

(91,) (91, 8)
```

```
# Importing and creating the DecisionTreemodel

from sklearn.tree import DecisionTreeRegressor
dtree = DecisionTreeRegressor(random_state=0)
dtree.fit(X_train, y_train)
```

DecisionTreeRegressor(random_state=0)

```
pred=dtree.predict(X)
y_pred_prob_df = pd.DataFrame(data=pred, columns=['Close_predict'])
df['Close_predict'] = y_pred_prob_df
df
```

/usr/local/lib/python3.8/dist-packages/sklearn/base.py:443: UserWarning:
X has feature names, but DecisionTreeRegressor was fitted without feature names

| | Date | Open | High | Low | Close | Adj Close | Volume | Close_predict |
|-----|------------|------------|------------|------------|------------|------------|----------|---------------|
| 0 | 2020-12-11 | 146.550003 | 151.500000 | 135.100006 | 139.250000 | 139.250000 | 26980800 | 139.250000 |
| 1 | 2020-12-14 | 135.000000 | 135.300003 | 125.160004 | 130.000000 | 130.000000 | 16966100 | 131.589996 |
| 2 | 2020-12-15 | 126.690002 | 127.599998 | 121.500000 | 124.800003 | 124.800003 | 10914400 | 124.800003 |
| 3 | 2020-12-16 | 125.830002 | 142.000000 | 124.910004 | 137.990005 | 137.990005 | 20409600 | 137.990005 |
| 4 | 2020-12-17 | 143.000000 | 152.449997 | 142.669998 | 147.050003 | 147.050003 | 15054700 | 147.050003 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 449 | 2022-09-26 | 101.779999 | 105.360001 | 101.559998 | 103.230003 | 103.230003 | 5278800 | 103.230003 |
| 450 | 2022-09-27 | 106.040001 | 108.169998 | 104.139999 | 106.370003 | 106.370003 | 5081000 | 106.370003 |
| 451 | 2022-09-28 | 106.700996 | 111.129997 | 105.779999 | 110.690002 | 110.690002 | 4518200 | 110.690002 |
| 452 | 2022-09-29 | 109.190002 | 109.474998 | 105.169998 | 106.660004 | 106.660004 | 4259500 | 106.660004 |



4532022-09-30104.760002109.080002104.730003105.040001105.0400014747800104.949997

454 rows × 8 columns

Decision Tree Regressor

```
from sklearn import metrics
import math
y_pred = dtree.predict(X_test)
mae = metrics.mean_absolute_error(y_test, y_pred)
mse = metrics.mean_squared_error(y_test, y_pred)
r2 = metrics.r2_score(y_test, y_pred)
rmse = math.sqrt(mse)
# accu = metrics.accuracy_score(y_test, y_pred)

print('MAE is {}'.format(mae))
print('MSE is {}'.format(mse))
print('R2 score is {}'.format(r2))
print('RMSE score is {}'.format(rmse))
# print('Accuracy score is {}'.format(accu*100))
```

MAE is 0.3470327142857169
MSE is 0.269499517831574
R2 score is 0.9996783892298543
RMSE score is 0.5191334296995079

Random Forest Regressor

```
from sklearn.ensemble import RandomForestRegressor
rf = RandomForestRegressor(random_state=0)
rf.fit(X_train, y_train)
```

RandomForestRegressor(random_state=0)

```
from sklearn import metrics
import math
y_pred = rf.predict(X_test)
mae = metrics.mean_absolute_error(y_test, y_pred)
mse = metrics.mean_squared_error(y_test, y_pred)
r2 = metrics.r2_score(y_test, y_pred)
rmse = math.sqrt(mse)

print('MAE is {}'.format(mae))
print('MSE is {}'.format(mse))
print('R2 score is {}'.format(r2))
print('RMSE score is {}'.format(rmse))
```

MAE is 0.29589454329670223
MSE is 0.2637844988276057
R2 score is 0.9996852093224394
RMSE score is 0.5135995510391396

Visualize Random Forest Regressor

```
pred_rdt=rf.predict(X)
y_pred_prob_df_rdt = pd.DataFrame(data=pred_rdt, columns=['Close_predict'])
df['Close_predict'] = y_pred_prob_df_rdt
df
```

/usr/local/lib/python3.8/dist-packages/sklearn/base.py:443: UserWarning:

X has feature names, but RandomForestRegressor was fitted without feature names

| | Date | Open | High | Low | Close | Adj Close | Volume | Close_predict |
|---|------------|------------|------------|------------|------------|------------|----------|---------------|
| 0 | 2020-12-11 | 146.550003 | 151.500000 | 135.100006 | 139.250000 | 139.250000 | 26980800 | 139.164899 |




| | | | | | | | | |
|-----|------------|------------|------------|------------|------------|------------|----------|------------|
| 1 | 2020-12-14 | 135.000000 | 135.300003 | 125.160004 | 130.000000 | 130.000000 | 16966100 | 130.719802 |
| 2 | 2020-12-15 | 126.690002 | 127.599998 | 121.500000 | 124.800003 | 124.800003 | 10914400 | 125.152402 |
| 3 | 2020-12-16 | 125.830002 | 142.000000 | 124.910004 | 137.990005 | 137.990005 | 20409600 | 136.527903 |
| 4 | 2020-12-17 | 143.000000 | 152.449997 | 142.669998 | 147.050003 | 147.050003 | 15054700 | 146.952300 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 449 | 2022-09-26 | 101.779999 | 105.360001 | 101.559998 | 103.230003 | 103.230003 | 5278800 | 103.366302 |
| 450 | 2022-09-27 | 106.040001 | 108.169998 | 104.139999 | 106.370003 | 106.370003 | 5081000 | 106.500002 |
| 451 | 2022-09-28 | 106.700996 | 111.129997 | 105.779999 | 110.690002 | 110.690002 | 4518200 | 110.838201 |
| 452 | 2022-09-29 | 109.190002 | 109.474998 | 105.169998 | 106.660004 | 106.660004 | 4259500 | 106.787003 |
| 453 | 2022-09-30 | 104.760002 | 109.080002 | 104.730003 | 105.040001 | 105.040001 | 4747800 | 105.485001 |

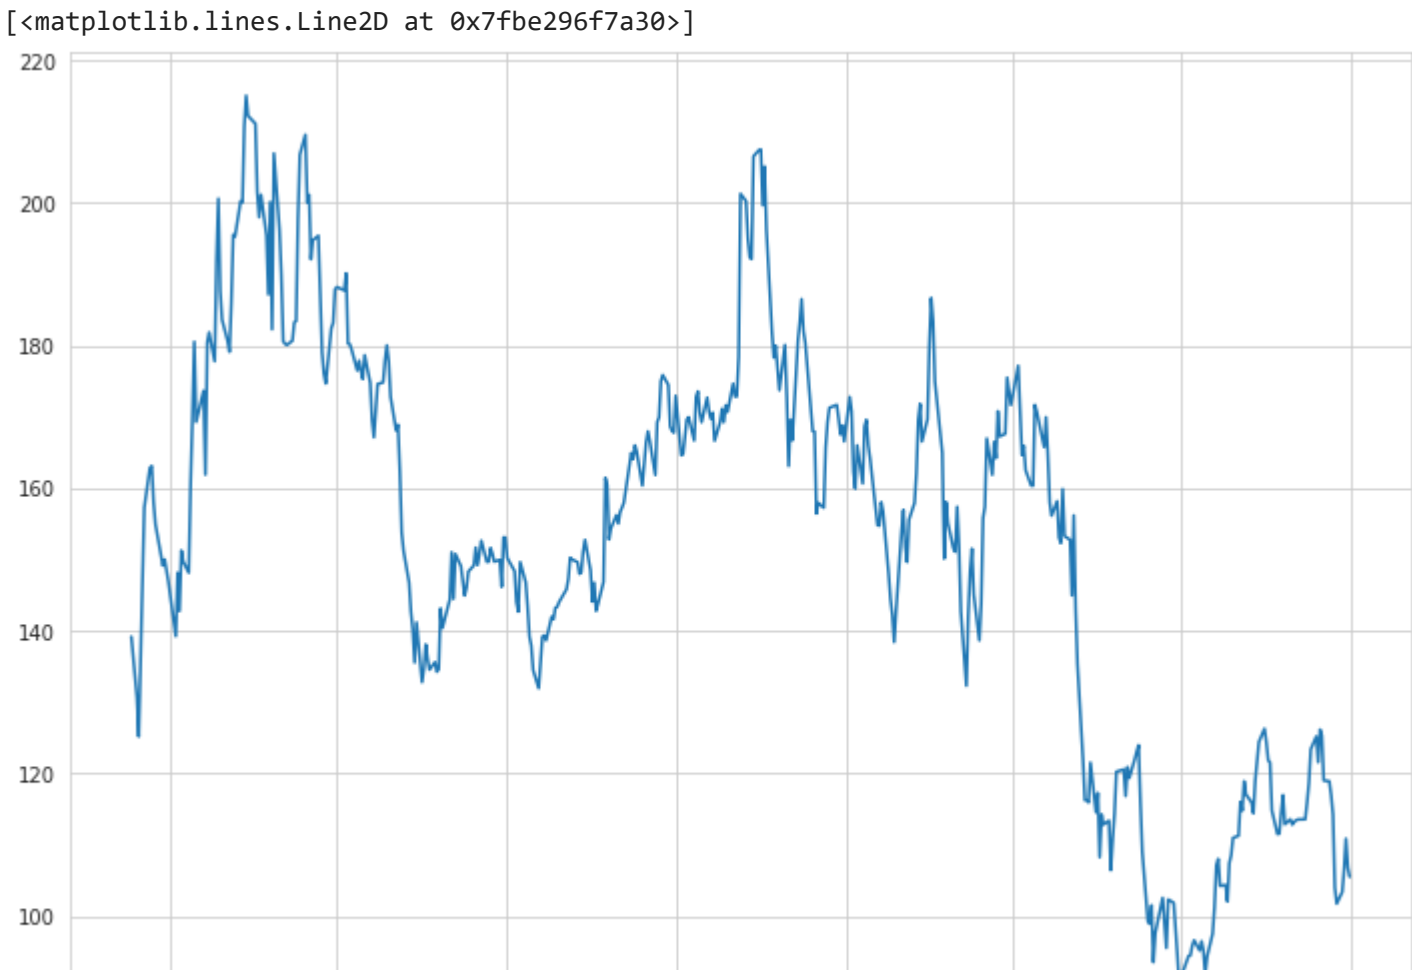
454 rows × 8 columns

```
Date = df['Date']
Close = df['Close']
Close2 = df['Close_predict']
```

```
plt.plot (Date,Close, label='Close')
plt.plot (Date,Close2, label='Close Predicted')
plt.title ('AirBnB Stock Prediction')
plt.xlabel ('Date')
plt.ylabel ('Close Price')
plt.legend(loc = 'upper right')
plt.xticks(rotation=90)
plt.show()
```



```
plt.figure(figsize=(12,9))
plt.subplot(1,1,1)
plt.plot(Date,Close2,label='Close Predicted')
```



| | | | | | | | |
|---------|---------|---------|---------|---------|---------|---------|---------|
| | | | | | | | |
| 2021-01 | 2021-04 | 2021-07 | 2021-10 | 2022-01 | 2022-04 | 2022-07 | 2022-10 |

[Colab paid products](#) - [Cancel contracts here](#)